Group 32 Progress Report: Predicting Purchase Outcomes from Online Shopping Session Behavior

Ankur Pandey, Allen Tong, Krish Haryani
{pandea23,tonga17,haryanik}@mcmaster.ca

1 Introduction

Most visits to an online store end without a purchase. If a retailer could estimate, from the way a session unfolds, how likely it is to end in a purchase, it could learn which browsing patterns are associated with conversion and where its website could be improved. Our project builds a machine learning application that makes this estimate for individual e-commerce sessions.

We frame the task as supervised, single-label binary classification on tabular data. Each record in the UCI Online Shoppers Purchasing Intention Dataset (Sakar and Kastro, 2018) summarizes one session with 17 inputs: page counts and time spent on administrative, informational and product pages; bounce and exit rates; the PageValues analytics measure; calendar context; and device, region, traffic-source and visitor-type categories. The target, `Revenue`, records whether the session ended in a purchase. The model estimates a purchase probability and assigns a class using a probability threshold selected on validation data.

Three properties make the problem harder than it first appears. First, purchases are rare: only 15.5% of the 12,330 sessions end in a purchase, so a model that always predicts "no purchase" is 84.5% accurate while being useless. Second, shopper behavior is heterogeneous, and the activity measures are long-tailed. Third, a single analytics feature, PageValues, carries much of the signal. In our training split, 55.6% of sessions with a positive PageValues end in a purchase, compared with 3.9% of sessions where it is zero. However, PageValues is derived from the value of pages that previously led to transactions. Not every website records it, sites may compute it differently, and it is unclear whether it would be available before a purchase is made. A model that relies on it may therefore be less reliable outside this dataset.

Our contribution is not a new leaderboard result on this dataset. Instead, we aim to provide (i) a careful comparison of majority-class baselines, regularized logistic regression, random forest and XGBoost under a fixed, leakage-aware evaluation protocol; (ii) a feature-availability study that compares all 17 inputs with subsets that exclude PageValues and, optionally, the bounce and exit rates; (iii) calibrated purchase probabilities rather than labels alone; and (iv) an inspectable Streamlit demonstration that returns probabilities, predicted classes and influential inputs for a session summary. Our results describe session-summary classification on one website and do not establish generalization to future periods or other stores.

2 Related Work

Purchase intent on this dataset. The dataset was introduced by Sakar et al. (2019), who proposed a real-time system with two modules. The first module predicts purchasing intention from session-level aggregate features like ours, comparing a multilayer perceptron with other classifiers such as random forests and support vector machines, and using oversampling and feature selection to address class imbalance. The second module uses an LSTM recurrent network on sequential clickstream data to estimate the likelihood that a visitor will leave the site. Our work extends their session-level study in two directions: we explicitly measure how much performance depends on PageValues and other analytics-derived inputs, and we evaluate probability quality in addition to classification accuracy.

Clickstream models of conversion. Earlier marketing research modeled conversion directly from browsing behavior. Moe and Fader (2004) modeled each visitor's probability of buying as it evolves across repeated visits, accounting for differences between shoppers, and showed that conversion behavior changes over time. Montgomery et al.

(2004) modeled the sequence of pages a user views and showed that the browsing path carries information that improves predictions of purchase conversion. Our dataset contains only per-session summaries without raw page sequences, so we cannot use path models directly, but these studies motivate the page-count and duration inputs we use.

Models for tabular data. Our candidate models are standard strong performers on structured data. Random forests (Breiman, 2001) average many decorrelated decision trees, and XGBoost (Chen and Guestrin, 2016) is a regularized, scalable implementation of gradient-boosted trees. A comparison by Shwartz-Ziv and Armon (2022) found that XGBoost generally matched or outperformed recent deep learning models on tabular benchmarks while requiring less tuning, which supports our choice to focus on tree ensembles and a logistic regression baseline instead of a neural network.

Imbalance, evaluation and calibration. He and Garcia (2009) survey learning from imbalanced data, including resampling and cost-sensitive methods, and note that overall accuracy can be misleading when one class dominates. Following Saito and Rehmsmeier (2015), who show that precision–recall analysis is more informative than ROC analysis when positives are rare, we use purchase-class average precision as our primary metric. Niculescu-Mizil and Caruana (2005) show that common learners such as boosted trees and random forests can produce systematically distorted probabilities that calibration methods correct, which motivates our use of Brier scores and reliability plots. Finally, Kaufman et al. (2012) describe how features that would not be available at prediction time can inflate reported performance. This concern underlies our PageValues study and our decision to fit all preprocessing within training folds only.

3 Dataset

3.1 Source and structure

We use the complete UCI Online Shoppers Purchasing Intention Dataset (Sakar and Kastro, 2018), introduced by Sakar et al. (2019) and distributed as a single CSV file under a CC BY 4.0 license. It contains 12,330 sessions, each from a different user over a one-year period, with 17 input features and the Boolean target `Revenue`, which records whether the session ended in a purchase (not the amount spent). No scraping, API access or manual annotation was required, since every session already carries its recorded outcome. Table 1 summarizes the inputs.

Inputs	Type
Administrative, Informational, ProductRelated (pages visited)	Count
The three matching `_Duration` fields (total time on each page type)	Continuous, units unspecified
BounceRates, ExitRates	Fraction in $[0, 1]$
PageValues (average value of pages visited before a transaction)	Continuous, currency unspecified
SpecialDay (closeness to a holiday)	Score in $\{0, 0.2, \ldots, 1\}$
Month, VisitorType	Nominal label
OperatingSystems, Browser, Region, TrafficType	Nominal integer code
Weekend	Boolean

Table 1: The 17 session-level inputs. Category codes have no documented meaning or order.

3.2 Properties of the data

Before dividing the file, we audited it in its entirety. Every numeric field is non-negative, and there are no missing or blank values. Five characteristics of the data influenced our design.

Class imbalance. Only 1,908 sessions (15.5%) end in a purchase, against 10,422 non-purchases. A classifier that never predicts a purchase is 84.5% accurate, which is why we do not use accuracy as our main metric (Section 6).

Skewed activity and meaningful zeros. All ten numeric inputs are positively skewed. Product-page time has a median of 598.9 but a maximum of 63,974, and sessions view a median of 18 product pages but only one administrative page and no informational pages. Zeros are common and informative rather than missing: PageValues is zero for 77.9% of sessions, and SpecialDay for 89.9%. Bounce and exit rates both stop at 0.2, with about 700 sessions each at exactly that value, which suggests a boundary in how the source computed them.

Correlated inputs. On the training split, each page count is strongly rank-correlated with its duration (Spearman ρ of 0.88–0.95), and bounce and exit rates are moderately correlated ($\rho = 0.60$). PageValues has by far the strongest association with purchase ($\rho = 0.62$), as discussed in the introduction.

Uneven categories. January and April are absent entirely, and the sessions per month range from 184

(February) to 3,364 (May). Purchase rates also vary by group (Figure 1): 25.5% of training sessions in November end in a purchase compared with 2.1% in February, and new visitors purchase more often than returning visitors (23.9% vs. 14.1%). Several browser, operating system and traffic-type codes occur fewer than 20 times, and some occur only once.

Identical sessions. 201 sessions fall into 76 groups whose 17 inputs are identical, and no group has conflicting labels. Because the data has no user identifiers, we cannot tell whether these are true duplicates or different users with the same summary, so we keep them all.

3.3 Preprocessing

We left the source file unchanged and removed no sessions, outliers or repeated rows; extreme values are kept unless there is evidence they are errors. All preprocessing decisions were made using the training split only.

Splitting. We divided the sessions into training, validation and test sets of 8,632, 1,849 and 1,849 sessions (about 70/15/15) with a fixed random seed. The split is stratified, so each set keeps the 15.5% purchase rate, and grouped, so identical sessions always land in the same set. This prevents a model from being evaluated on an exact copy of a training example. The training set is further divided into five fixed cross-validation folds under the same rules. The assignments are saved to files and verified on every load, so every team member uses the same split (Section 6).

Input checks and encoding. Every session passes through an input check that rejects missing values, negative counts, rates outside $[0, 1]$ and invalid category codes, without silently filling in or clipping values. The six nominal fields are one-hot encoded, `Weekend` becomes a 0/1 indicator, and numeric inputs are standardized for logistic regression but left unscaled for tree-based models. These steps are combined with each model in a single scikit-learn pipeline, so the encoder's categories and the scaler's statistics are learned from the training rows of each fold only. Categories that were not seen during training are encoded as all zeros. With all 17 inputs, this produces 74 model features (Section 4).

Limitations. The released data contains session summaries only, with no timestamps, year, user identifiers or page-by-page click sequences. We therefore cannot build a time-ordered split or show that a model would generalize to later periods or other websites.

Input block	Columns
Numeric inputs	10
Month	10
OperatingSystems	8
Browser	13
Region	9
TrafficType	20
VisitorType	3
Weekend	1
Total	**74**

Table 2: Encoded model features for all 17 inputs, with category levels counted on the training split.

4 Features

Our models receive the 17 session-level inputs in Table 1, encoded as described in Section 3.3. Fitting the encoder on the training split gives 74 model features (Table 2). Each nominal field is represented by one indicator per observed level instead of its integer code, since the codes have no documented order.

Feature engineering. We have not created any extra features so far. There are no interaction terms, ratios or polynomial expansions, and the skewed numeric inputs enter the models without a log transform. We also do not use neural networks or learned embeddings, so each model receives a fixed tabular vector.

Augmentation. We did not augment the data. No sessions were duplicated or synthesized, and unlike Sakar et al. (2019), we did not oversample the purchase class. Class imbalance is instead handled in the loss through class weights (Section 5).

Feature sets. We do not use automatic feature selection. Instead, we defined three feature sets in advance to measure how much performance depends on inputs that may be hard to reproduce at prediction time. The first uses all 17 inputs (74 features). The second removes PageValues (16 inputs, 73 features), since it summarizes past page values and the measurements it needs may not exist before a purchase. The third also removes BounceRates and ExitRates (14 inputs, 71 features), which depend on page-level analytics definitions and history. The data does not say when each measurement was made, so these sets are not claims of leakage or of

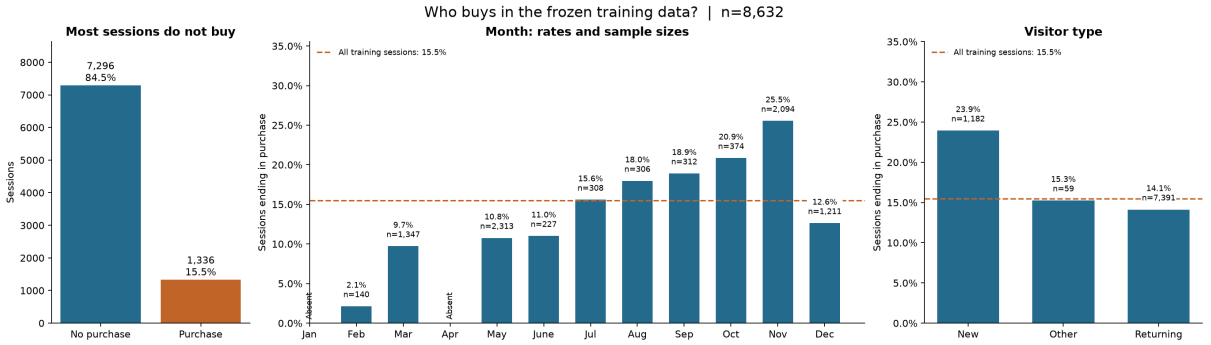


Figure 1: Training split ($n = 8{,}632$). Left: class balance. Centre and right: share of sessions ending in a purchase by month and by visitor type, with sample sizes. The dashed line marks the overall purchase rate (15.5%).

deployment readiness.

All three sets use the same data splits and cross-validation folds, and each is fitted as its own pipeline, so a removed input is truly absent rather than filled with zero at prediction time. Our current logistic regression model uses the full 17-input set (Section 6); the comparison across feature sets is planned for the next stage of the project.

5 Implementation

Our implementation uses Python, pandas and NumPy for data handling, scikit-learn for preprocessing and classification, and matplotlib for the evaluation figures. Each classifier is combined with the preprocessing steps in one pipeline. This lets us pass raw session records to the fitted model without having to encode or scale them again by hand.

Baselines. We implemented two simple baselines before training a classifier. The majority-class baseline always predicts no purchase. The class-prior baseline gives every session the same purchase probability, equal to the purchase rate in the training set (0.1548). Neither uses browsing behavior, so they establish what can be achieved without learning from the session inputs.

Logistic regression. Our first trained classifier is L2-regularized logistic regression on all 17 inputs. For a session with encoded features $\mathbf{x}_i$ and label $y_i \in \{0, 1\}$, the model predicts $p_i = \sigma(\mathbf{w}^\top \mathbf{x}_i + b)$. Training minimizes

$$\frac{1}{2}\|\mathbf{w}\|^2 + C \sum_{i=1}^{n} s_i \, \ell_i,$$

where $\ell_i = -y_i \log p_i - (1 - y_i) \log(1 - p_i)$ is the log-loss and s_i is a class weight. The log-loss penalizes confident wrong probabilities, and the first term discourages large coefficients. C is the inverse regularization strength, so smaller values apply a stronger penalty. In an unweighted fit, every $s_i = 1$. Balanced weights are inversely proportional to class frequency, so a purchase counts about 5.5 times as much as a non-purchase, without adding or removing training sessions.

Optimization. We minimize this objective with L-BFGS, a quasi-Newton method. Like gradient descent, it follows the gradient of the loss, but it also estimates the curvature from its recent steps, which lets it take better-scaled steps and converge in fewer iterations. The objective is convex and the training set is small, so full-batch optimization is cheap and we did not need stochastic gradient descent. We used a tolerance of 10^{-4} and a maximum of 2,000 iterations. We compared four values of C, $\{0.01, 0.1, 1, 10\}$, each with and without balanced class weights.

These eight settings were compared on the five fixed training folds. The encoder, scaler and classifier were fitted separately within each fold, and we selected the setting with the highest mean purchase average precision. The selected model used $C = 0.1$ with balanced class weights and was then refitted on all 8,632 training sessions. All 40 cross-validation fits and the final fit converged, and the final fit took 44 iterations.

Reproducibility. We saved the complete fitted pipeline with joblib, reloaded it and verified that its predictions matched exactly on all 1,849 validation sessions. The evaluation code exports scores, predicted classes, metrics and confusion matrices with the split identifier and seed. Random forest, XGBoost and comparisons across the reduced feature sets have not yet been evaluated; they form the next stage of the project.

Model	Threshold	AP	ROC-AUC	Precision	Recall	F1	Brier
Majority class	0.5000	0.1547	0.5000	0.0000	0.0000	0.0000	0.1547
Class prior	0.5000	0.1547	0.5000	0.0000	0.0000	0.0000	0.1308
Logistic regression	0.5000	0.6644	0.9008	0.5216	0.7587	0.6182	0.1197
Logistic regression	0.6057	0.6644	0.9008	0.6242	0.6678	0.6453	0.1197

Table 3: Validation results for the baselines and selected logistic regression model. The threshold of 0.6057 was selected on these same validation sessions, so its F1 is a development result. Higher is better for all reported metrics except Brier.

6 Results and Evaluation

Evaluation strategy. We use the fixed grouped, stratified split described in Section 3.3. Cross-validation within training selects model settings; validation supports development comparisons and threshold selection. All results below are from training cross-validation or the validation set, which contains 286 purchases among 1,849 sessions. We have not evaluated test performance. Before the final test evaluation, we will freeze the fitted models, feature sets, calibration methods and thresholds.

Our primary metric is purchase average precision (AP), which summarizes precision and recall as the score threshold changes. We also report ROC-AUC, purchase precision, recall and F1, a confusion matrix, and the Brier score and a reliability plot to examine probability quality. AP and ROC-AUC use the model scores rather than the predicted classes. For class predictions, we report the reference threshold of 0.5 and a threshold chosen to maximize purchase F1 on validation. We use F1 as a compromise between precision and recall because we do not have reliable business costs for either kind of error.

Preliminary results. Table 3 shows that logistic regression learns useful information beyond the class balance: its AP is 0.6644, compared with 0.1547 for both baselines. At threshold 0.5, it finds 217 of the 286 purchases, with 199 false alarms. Raising the threshold to 0.6057 reduces false alarms to 115 but finds only 191 purchases. Precision therefore rises from 52.2% to 62.4%, while recall falls from 75.9% to 66.8%. The resulting F1 of 0.6453 is optimistic because the same validation sessions were used to choose that threshold. For comparison, maximizing F1 for either constant-score baseline predicts every session as a purchase and gives an F1 of 0.2679.

Class weighting and probability quality. Table 4 compares all eight settings. They fall within

	Mean CV AP		Mean CV Brier	
C	None	Bal.	None	Bal.
0.01	0.6428	0.6430	0.0871	0.1264
0.1	0.6418	**0.6469**	0.0851	0.1212
1	0.6370	0.6447	0.0853	0.1214
10	0.6343	0.6430	0.0856	0.1218

Table 4: Logistic regression on the five training folds, without class weights (None) and with balanced weights (Bal.). The selected setting is in bold. Lower Brier is better.

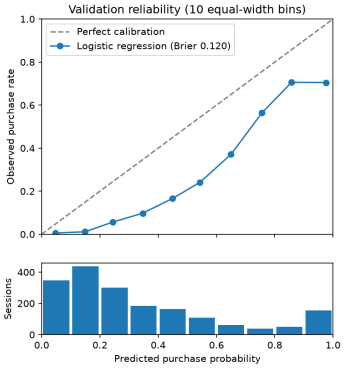


Figure 2: Validation reliability plot for the selected model. Points below the dashed line are overestimates. The lower panel counts sessions per bin.

0.013 AP of each other, well inside the fold-to-fold standard deviation of about 0.028, so weighting and the choice of C make no clear difference to ranking. Weighting does clearly hurt probability quality, raising mean Brier from about 0.085 to about 0.12. On validation, the selected model's Brier score is 0.1197, only modestly better than the class-prior baseline's 0.1308. The reliability plot (Figure 2) shows that the model overestimates purchase probability in every bin. For example, sessions in the 0.6–0.7 bin have an average predicted probability of 0.65 but purchase only 37% of the time. The model ranks sessions usefully, but its probabilities should not yet be read as purchase chances.

7 Feedback and Plans

During our consultation, we walked the TA through our preprocessing, five-fold cross-validation setup

and logistic regression results, and raised our concern about class imbalance.

The TA suggested that a random baseline can be a useful comparison in some cases, with performance above that baseline providing an initial check that the model is learning useful information. We currently compare against majority-class and class-prior baselines. The class-prior baseline gives every session the same score, so it does not make random guesses. To follow up on the advice, we plan to add a baseline that randomly predicts purchase according to the training purchase rate and repeat it with several seeds to show how much its results vary.

The TA also encouraged us to try models beyond logistic regression. We explained that random forest and XGBoost are already part of our plan and will be implemented next. Both can learn nonlinear relationships between the session inputs. We will compare them using the same splits, training folds and metrics, so any improvement can be assessed consistently.

What we will change. Our first model changed the plan in two ways. First, better ranking did not mean better probabilities. Class weighting gave no clear AP gain but worse Brier scores, so we will treat unweighted models as the default and compare them with models calibrated on out-of-fold training predictions. We will not present the current weighted probabilities as purchase chances. Second, the numeric inputs are heavily skewed but enter logistic regression unchanged. We will test a zero-preserving log transform within the training folds, while tree models keep the original values. We also still need to run the feature-set study with and without PageValues, and to inspect errors by month and visitor type.

Timeline. We plan to finish random forest, XGBoost, tuning, the feature-set study, calibration and error analysis by November 28. By December 3, we will freeze the selected model and threshold, run the single test evaluation and connect the saved pipeline to the Streamlit demonstration. The final week is reserved for the report and demo.

Team Contributions

Ankur worked on the project setup, dataset checks and exploratory analysis. He also prepared the data splits and preprocessing pipeline used for the model experiments.

Allen implemented the baseline models and logistic regression. He evaluated their performance and prepared the results and figures used to compare the models.

Krish organized the progress report and wrote the introduction and related work. He also prepared the dataset and feature descriptions, explaining the data properties and how the inputs are represented.

Generative AI Disclosure

We used generative AI throughout this project, following the course policy. As a team, we first discussed the overall goals and divided the work into smaller tasks. We then used AI to turn the notes from those discussions into the written GitHub issues we used to track the work. The decisions in that plan, including the dataset, the models to compare and the evaluation protocol, came from our own discussions.

Within individual tasks, we used AI for background research, for building intuition about the methods we applied, and for writing parts of the code, especially where we did not know the syntax of a library. We asked it to explain the code it produced, so we could understand each step and remove anything the task did not need, and we checked the results with our unit tests and by rerunning the pipeline. We also used AI to help draft and edit the text of this report. All reported numbers come from our saved result files. We reviewed the AI-generated code and text and take full responsibility for this submission.

Carbon footprint. We did not record compute time, hardware or region, so we used the course's estimate of 4.32 g of CO_2 per query. Across roughly 35 to 40 queries, this gives about 151 to 173 g of CO_2.

References

Leo Breiman. 2001. Random forests. *Machine Learning*, 45(1):5–32.

Tianqi Chen and Carlos Guestrin. 2016. XGBoost: A scalable tree boosting system. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 785–794.

Haibo He and Edwardo A. Garcia. 2009. Learning from imbalanced data. *IEEE Transactions on Knowledge and Data Engineering*, 21(9):1263–1284.

Shachar Kaufman, Saharon Rosset, Claudia Perlich, and Ori Stitelman. 2012. Leakage in data mining: Formulation, detection, and avoidance. *ACM Transactions on Knowledge Discovery from Data*, 6(4):1–21.

Wendy W. Moe and Peter S. Fader. 2004. Dynamic conversion behavior at e-commerce sites. *Management Science*, 50(3):326–335.

Alan L. Montgomery, Shibo Li, Kannan Srinivasan, and John C. Liechty. 2004. Modeling online browsing and path analysis using clickstream data. *Marketing Science*, 23(4):579–595.

Alexandru Niculescu-Mizil and Rich Caruana. 2005. Predicting good probabilities with supervised learning. In *Proceedings of the 22nd International Conference on Machine Learning*, pages 625–632.

Takaya Saito and Marc Rehmsmeier. 2015. The precision-recall plot is more informative than the ROC plot when evaluating binary classifiers on imbalanced datasets. *PLOS ONE*, 10(3):e0118432.

C. Okan Sakar and Yomi Kastro. 2018. Online shoppers purchasing intention dataset. UCI Machine Learning Repository. Licensed under CC BY 4.0.

C. Okan Sakar, S. Olcay Polat, Mete Katircioglu, and Yomi Kastro. 2019. Real-time prediction of online shoppers' purchasing intention using multilayer perceptron and LSTM recurrent neural networks. *Neural Computing and Applications*, 31(10):6893–6908.

Ravid Shwartz-Ziv and Amitai Armon. 2022. Tabular data: Deep learning is not all you need. *Information Fusion*, 81:84–90.